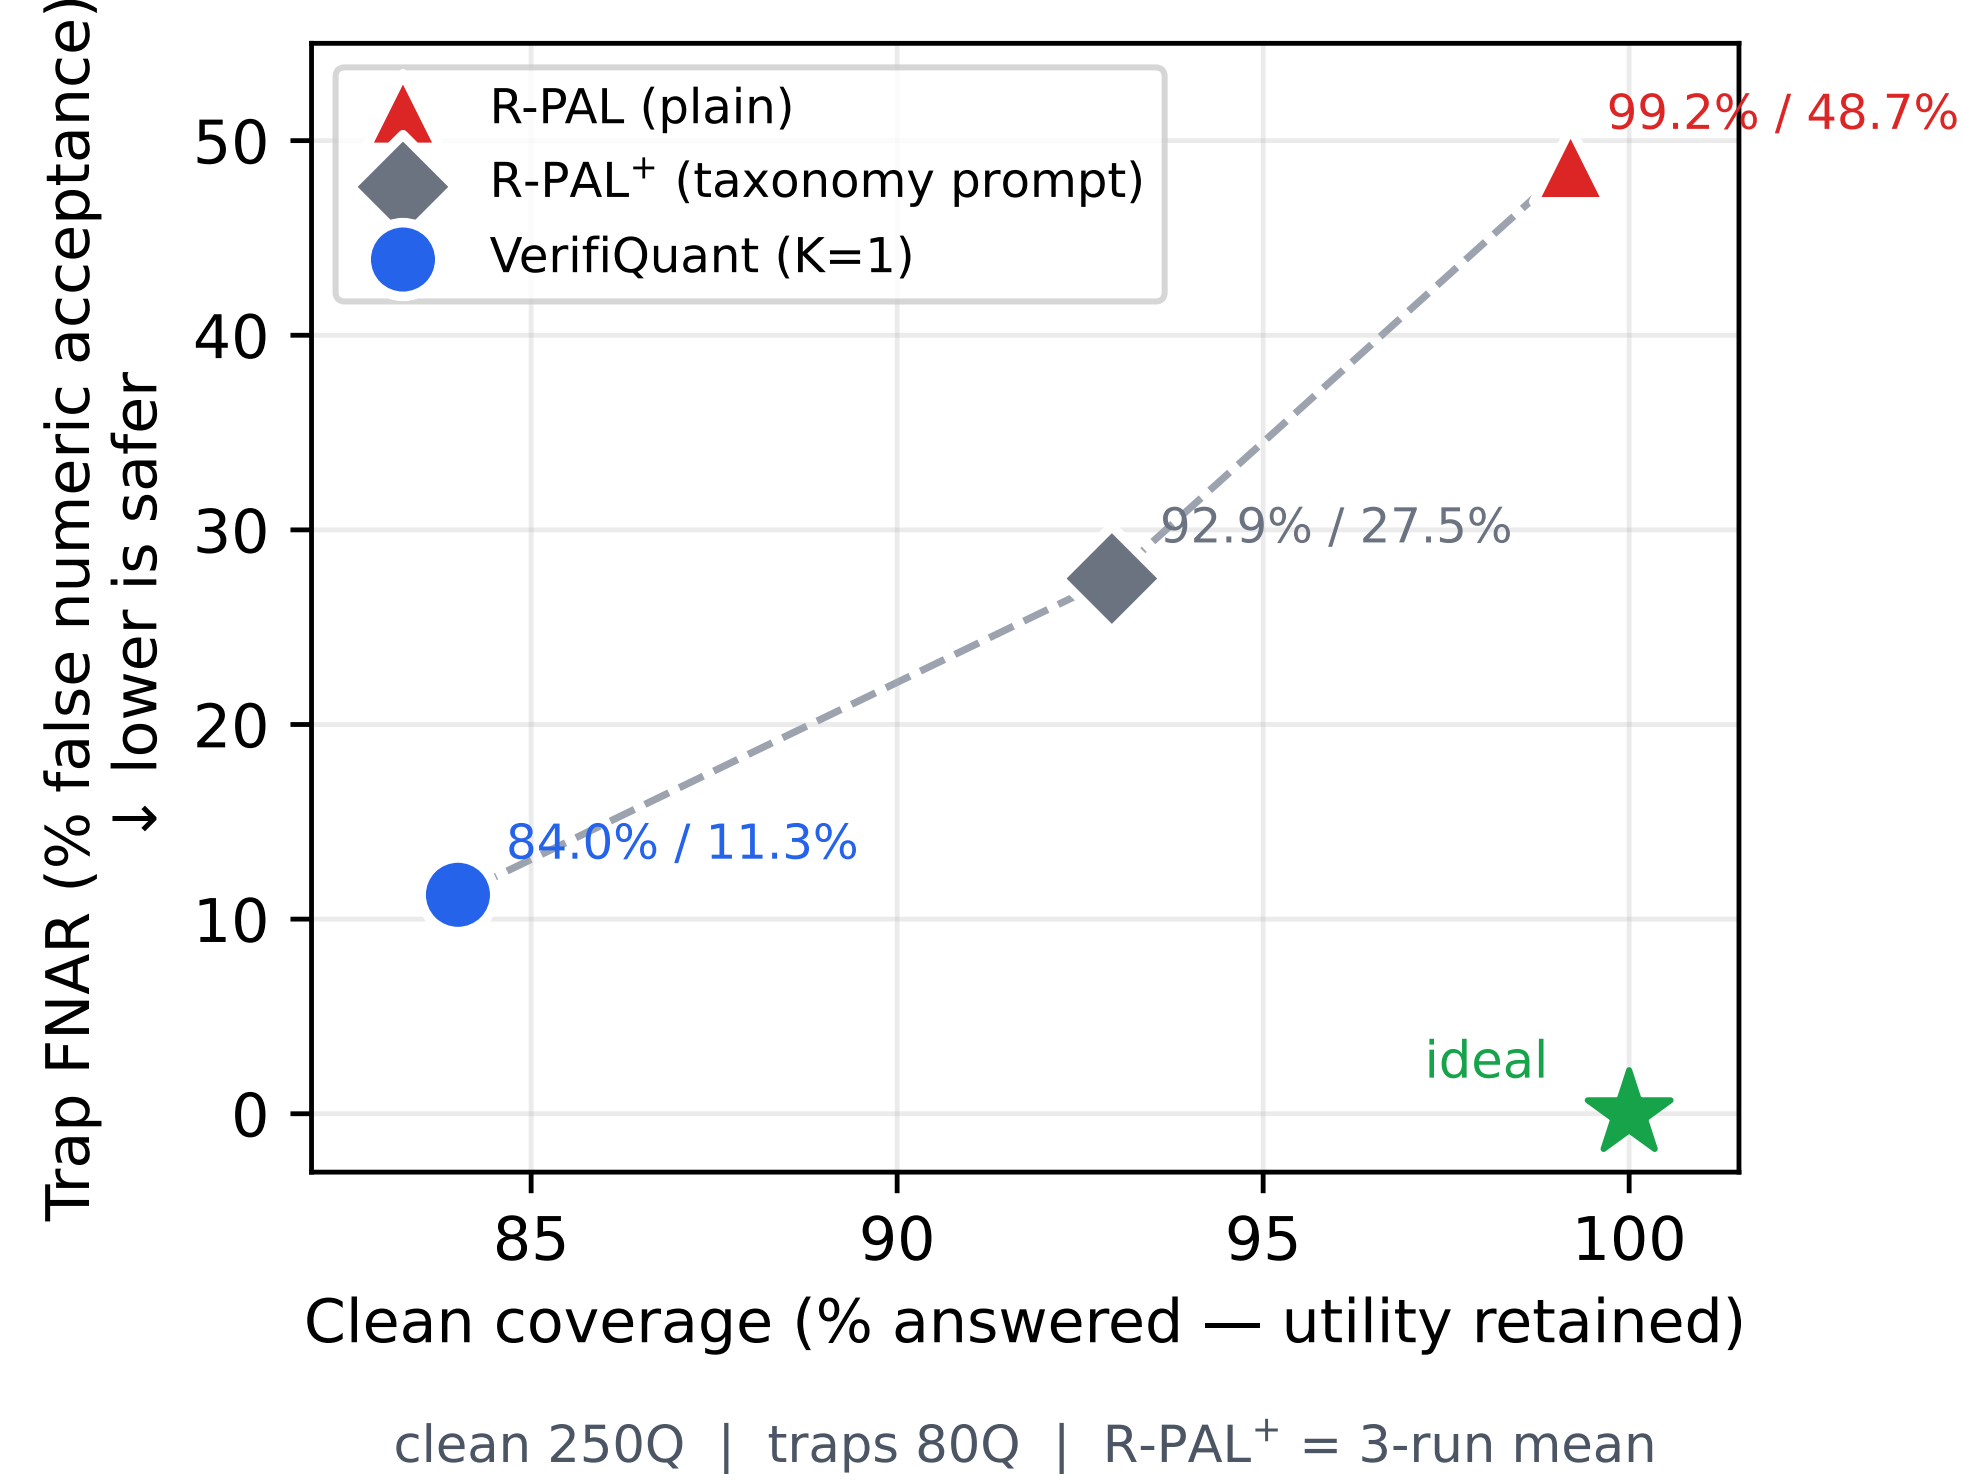
R-PAL (plain)
R-PAL+ (taxonomy prompt)
VerifiQuant (K=1)
99.2% / 48.7%
92.9% / 27.5%
84.0% / 11.3%
ideal
Trap FNAR (% false numeric acceptance)
↓ lower is safer
Clean coverage (% answered — utility retained)
0
10
20
30
40
50
85
90
95
100
clean 250Q | traps 80Q | R-PAL+ = 3-run mean